

Data Preprocessing:

The data has been uploaded and pre-processed to ensure that it is cleaned, formatted and outliers are addressed before exploring relationships between variables and analyzing correlations.

1. **Column Names:** Rename columns to have clear, short and lowercase names.
2. **Data Types:** assign / Convert the data type of each column to appropriate ones.
3. **Duplicate Rows:** Identify and drop duplicate rows. 1635 duplicates found
4. **Missing Values:** Check for missing values and handle them appropriately. No missing value found
5. **Aggregation of feature categories** to fewer levels: Age, Education and Income
6. **Frequency distribution** of the target/outcome variable, diabetes
7. **Outliers:** Detect, visualize and remove outliers.
8. **Histograms:** Plot histograms to visualize the distribution of numeric features.
9. **Scatter Plot/Bar Plot:** Plot scatter plots or bar plots to explore relationships between variables.
10. **Summary Statistics:** Calculate summary statistics for numeric features.
11. **Correlation Matrix/Heat Map:** Calculate and visualize the correlation between numeric features
12. **Chi-square test** for independence of categorical variables against 'diabetes'
13. **The Frequency distribution** of some categorical variables: 'age', 'education', 'income'
14. **The distribution of 'diabetes' outcome** by risk factors
15. **The distribution of diabetes outcome** by good habit features
16. **Prevalence of Diabetes** by age, gender, education and income level

Import necessary libraries

```
In [ ]: ## IMPORT NECESSARY PYTHON LIBRARIES

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn
import scipy.stats as stats
```

Read the diabetes dataset

```
In [ ]: df_diabetes = pd.read_csv("C:/Users/yitay/Documents/CIND820_BigDataAnalyticsProject

In [ ]: # see the first 5 observations
df_diabetes.head()
```

```
Out[ ]:
```

	Diabetes_binary	HighBP	HighChol	CholCheck	BMI	Smoker	Stroke	HeartDiseaseorA
0	0.0	1.0	0.0	1.0	26.0	0.0	0.0	
1	0.0	1.0	1.0	1.0	26.0	1.0	1.0	
2	0.0	0.0	0.0	1.0	26.0	0.0	0.0	
3	0.0	1.0	1.0	1.0	28.0	1.0	0.0	
4	0.0	0.0	0.0	1.0	29.0	1.0	0.0	

5 rows × 22 columns

1.Column Names: renamed columns to have clear, short and lowercase names.

```
In [ ]: # Lowercase, shorter, clear names
new_column_names = ['diabetes', 'bp', 'chol', 'cholcheck', 'bmi', 'smoker', 'stroke',
                    'education', 'income']
df_diabetes.columns = new_column_names
```

```
In [ ]: # Check the new column names
df_diabetes.columns
```

```
Out[ ]: Index(['diabetes', 'bp', 'chol', 'cholcheck', 'bmi', 'smoker', 'stroke',
              'heart_disease', 'activity', 'fruits', 'veggies', 'alcohol',
              'healthcare', 'nodocbccost', 'genhlth', 'menthlth', 'phyhlth', 'walk',
              'sex', 'age', 'education', 'income'],
              dtype='object')
```

```
In [ ]: # the first 5 obsrvations/rows after renaming the col
df_diabetes.head()
```

```
Out[ ]:
```

	diabetes	bp	chol	cholcheck	bmi	smoker	stroke	heart_disease	activity	fruits	...
0	0.0	1.0	0.0	1.0	26.0	0.0	0.0	0.0	1.0	0.0	...
1	0.0	1.0	1.0	1.0	26.0	1.0	1.0	0.0	0.0	1.0	...
2	0.0	0.0	0.0	1.0	26.0	0.0	0.0	0.0	1.0	1.0	...
3	0.0	1.0	1.0	1.0	28.0	1.0	0.0	0.0	1.0	1.0	...
4	0.0	0.0	0.0	1.0	29.0	1.0	0.0	0.0	1.0	1.0	...

5 rows × 22 columns

2. Data Types and structure: check and assign appropriate data types to each attribute

Check data types of attributes

```
In [ ]: # check the data type of attributes
print("Data Types of Attributes:")
df_diabetes.dtypes
```

Data Types of Attributes:

```
Out[ ]: diabetes      float64
bp                float64
chol              float64
cholcheck         float64
bmi               float64
smoker            float64
stroke            float64
heart_disease     float64
activity          float64
fruits            float64
veggies           float64
alcohol           float64
healthcare        float64
nodocbccost       float64
genhlth           float64
menthlth          float64
phyhlth           float64
walk              float64
sex               float64
age               float64
education         float64
income            float64
dtype: object
```

Check the data structure (rows and columns)

```
In [ ]: # (row, columns)
print(df_diabetes.shape)
print("Number of rows:", df_diabetes.shape[0])    # 0 for rows
print("Number of columns:", df_diabetes.shape[1])  # 1 for columns
```

(70692, 22)

Number of rows: 70692

Number of columns: 22

Count the number of unique values for each column

```
In [ ]: # Dictionary to store the count of unique values for each column
unique_counts = {}

# Loop through each column and get the count of unique values
for column in df_diabetes.columns:
    unique_counts[column] = df_diabetes[column].nunique()

# Print the number of unique values for each column
for column, count in unique_counts.items():
    print(f"Number of unique values - '{column}': {count}")
```

```
Number of unique values - 'diabetes': 2
Number of unique values - 'bp': 2
Number of unique values - 'chol': 2
Number of unique values - 'cholcheck': 2
Number of unique values - 'bmi': 80
Number of unique values - 'smoker': 2
Number of unique values - 'stroke': 2
Number of unique values - 'heart_disease': 2
Number of unique values - 'activity': 2
Number of unique values - 'fruits': 2
Number of unique values - 'veggies': 2
Number of unique values - 'alcohol': 2
Number of unique values - 'healthcare': 2
Number of unique values - 'nodocbccost': 2
Number of unique values - 'genhlth': 5
Number of unique values - 'menthlth': 31
Number of unique values - 'phyhlth': 31
Number of unique values - 'walk': 2
Number of unique values - 'sex': 2
Number of unique values - 'age': 13
Number of unique values - 'education': 6
Number of unique values - 'income': 8
```

Convert the data type of categorical features

```
In [ ]: # List of columns to convert to integer dtype
columns_to_integer = ['diabetes', 'bp', 'chol', 'cholcheck', 'smoker', 'stroke', 'h

# Convert the specified columns to integer dtype
for column in columns_to_integer:
    # Fill NaN values with 0 and convert to integer
    df_diabetes[column] = df_diabetes[column].astype(int)

# Check the data types
print(df_diabetes.dtypes)

# Print the DataFrame
print(df_diabetes)
```

```

diabetes      int32
bp            int32
chol          int32
cholcheck     int32
bmi           float64
smoker        int32
stroke        int32
heart_disease int32
activity      int32
fruits        int32
veggies       int32
alcohol       int32
healthcare    int32
nodocbccost   int32
genhlth       int32
menthlth      float64
phyhlth       float64
walk          int32
sex           int32
age           int32
education     int32
income        int32

```

```
dtype: object
```

```

      diabetes  bp  chol  cholcheck  bmi  smoker  stroke  heart_disease  \
0             0  1    0           1  26.0      0      0             0
1             0  1    1           1  26.0      1      1             0
2             0  0    0           1  26.0      0      0             0
3             0  1    1           1  28.0      1      0             0
4             0  0    0           1  29.0      1      0             0
...          ...  ..   ...         ...   ...     ...     ...             ...
70687         1  0    1           1  37.0      0      0             0
70688         1  0    1           1  29.0      1      0             1
70689         1  1    1           1  25.0      0      0             1
70690         1  1    1           1  18.0      0      0             0
70691         1  1    1           1  25.0      0      0             1

```

```

      activity  fruits  ...  healthcare  nodocbccost  genhlth  menthlth  \
0             1      0  ...           1             0      3      5.0
1             0      1  ...           1             0      3      0.0
2             1      1  ...           1             0      1      0.0
3             1      1  ...           1             0      3      0.0
4             1      1  ...           1             0      2      0.0
...          ...     ...  ...         ...         ...     ...     ...
70687         0      0  ...           1             0      4      0.0
70688         0      1  ...           1             0      2      0.0
70689         0      1  ...           1             0      5     15.0
70690         0      0  ...           1             0      4      0.0
70691         1      1  ...           1             0      2      0.0

```

```

      phyhlth  walk  sex  age  education  income
0          30.0    0    1    4           6      8
1           0.0    0    1   12           6      8
2          10.0    0    1   13           6      8
3           3.0    0    1   11           6      8
4           0.0    0    0    8           5      8
...          ...    ...  ...   ...         ...     ...

```

70687	0.0	0	0	6	4	1
70688	0.0	1	1	10	3	6
70689	0.0	1	0	13	6	4
70690	0.0	1	0	11	2	4
70691	0.0	0	0	9	6	2

[70692 rows x 22 columns]

3. Duplicate Rows: Identify and drop duplicate rows

- 1635 duplicate rows found

```
In [ ]: # Check for duplicate rows
duplicate_rows = df_diabetes[df_diabetes.duplicated()]

# Print the duplicate rows
if not duplicate_rows.empty:
    print("Duplicate rows found:")
    print(duplicate_rows)
else:
    print("No duplicate rows found.")

# Drop duplicates, keeping the first occurrence
df_diabetes = df_diabetes.drop_duplicates(keep='first')

# Print the cleaned DataFrame
print("\nDataFrame after removing duplicates:")
print(df_diabetes)
```

Duplicate rows found:

	diabetes	bp	chol	cholcheck	bmi	smoker	stroke	heart_disease	\
602	0	0	0	1	22.0	0	0	0	
689	0	0	0	1	26.0	0	0	0	
891	0	0	0	1	24.0	0	0	0	
1092	0	0	0	1	21.0	0	0	0	
1326	0	1	0	1	29.0	0	0	0	
...	
69865	1	1	1	1	27.0	1	0	0	
69939	1	1	1	1	27.0	1	0	0	
70305	1	1	0	1	30.0	0	0	0	
70591	1	1	1	1	30.0	0	0	1	
70663	1	1	1	1	33.0	0	0	0	

	activity	fruits	...	healthcare	nodocbccost	genhlth	menthlth	\
602	1	1	...	1	0	1	0.0	
689	1	1	...	1	0	1	0.0	
891	1	1	...	1	0	1	0.0	
1092	1	1	...	1	0	1	0.0	
1326	1	0	...	1	0	2	0.0	
...	
69865	1	1	...	1	0	4	0.0	
69939	1	1	...	1	0	3	0.0	
70305	0	0	...	1	0	3	0.0	
70591	1	1	...	1	0	2	0.0	
70663	1	1	...	1	0	3	0.0	

	phyhlth	walk	sex	age	education	income
602	0.0	0	0	6	6	8
689	0.0	0	0	6	6	8
891	0.0	0	0	6	6	8
1092	0.0	0	0	5	6	8
1326	0.0	0	1	10	5	6
...
69865	0.0	0	1	12	6	8
69939	0.0	0	1	12	6	8
70305	0.0	0	1	9	4	7
70591	0.0	0	1	10	6	8
70663	0.0	0	1	9	6	6

[1635 rows x 22 columns]

DataFrame after removing duplicates:

	diabetes	bp	chol	cholcheck	bmi	smoker	stroke	heart_disease	\
0	0	1	0	1	26.0	0	0	0	
1	0	1	1	1	26.0	1	1	0	
2	0	0	0	1	26.0	0	0	0	
3	0	1	1	1	28.0	1	0	0	
4	0	0	0	1	29.0	1	0	0	
...	
70687	1	0	1	1	37.0	0	0	0	
70688	1	0	1	1	29.0	1	0	1	
70689	1	1	1	1	25.0	0	0	1	
70690	1	1	1	1	18.0	0	0	0	
70691	1	1	1	1	25.0	0	0	1	

	activity	fruits	...	healthcare	nodocbcost	genhlth	menthlth	\
0	1	0	...	1	0	3	5.0	
1	0	1	...	1	0	3	0.0	
2	1	1	...	1	0	1	0.0	
3	1	1	...	1	0	3	0.0	
4	1	1	...	1	0	2	0.0	
...	
70687	0	0	...	1	0	4	0.0	
70688	0	1	...	1	0	2	0.0	
70689	0	1	...	1	0	5	15.0	
70690	0	0	...	1	0	4	0.0	
70691	1	1	...	1	0	2	0.0	

	phyhlth	walk	sex	age	education	income
0	30.0	0	1	4	6	8
1	0.0	0	1	12	6	8
2	10.0	0	1	13	6	8
3	3.0	0	1	11	6	8
4	0.0	0	0	8	5	8
...
70687	0.0	0	0	6	4	1
70688	0.0	1	1	10	3	6
70689	0.0	1	0	13	6	4
70690	0.0	1	0	11	2	4
70691	0.0	0	0	9	6	2

[69057 rows x 22 columns]

4. Missing Values: Identify and handle missing values.

- No missing value found

```
In [ ]: # Count the number of missing values in each column
missing_count_per_column = df_diabetes.isnull().sum()
missing_count_per_column
```



```
Out[ ]: diabetes      0
        bp            0
        chol          0
        cholcheck     0
        bmi           0
        smoker        0
        stroke        0
        heart_disease 0
        activity      0
        fruits        0
        veggies       0
        alcohol       0
        healthcare    0
        nodulebccost  0
        genhlth       0
        menthlth      0
        phyhlth       0
        walk          0
        sex           0
        age           0
        education     0
        income        0
        dtype: int64
```

```
In [ ]: ## Count the total number of missing values in the entire DataFrame
        total_missing_count = df_diabetes.isnull().sum().sum()
        print("Total number of missing values:",total_missing_count)
```

Total number of missing values: 0

5. Aggregation of feature categories to fewer levels: Age, Education and Income

Age

The age category levels have been reduced from 13 to 4. Levels 1 and 2 assigned to level 0 (young adults); 3 to 6 to level 1 (middle aged), 7 to 9 to level 2 (older adults) and 10 to 13 to level 3 (elderly)

```
In [ ]: # Define age mapping: maps original age levels to new categorical levels
        age_mapping = {
            1: 0, # Levels 1 & 2 map to 0 (Young Adults)
            2: 0,
            3: 1, # Levels 3 - 6 map to 1 (Middle Aged)
            4: 1,
            5: 1,
            6: 1,
            7: 2, # Levels 7 - 9 map to 2 (Older Adults)
            8: 2,
            9: 2,
            10: 3, # Levels 10 - 13 map to 3 (Elderly)
            11: 3,
            12: 3,
            13: 3
        }
```

```

## Replace the age levels in the DataFrame
df_diabetes['age'] = df_diabetes['age'].map(age_mapping)

## Rename the age levels to Young Adults, Middle Aged, Older Adults, and Elderly
#df_diabetes['age'] = df_diabetes['age'].replace({0: 'Young Adults', 1: 'Middle Age

## We can skip the renaming step if we prefer using numerical values instead of st
## Now, 'age' column will have levels 'Young Adults', 'Middle Aged', 'Older Adults'

```

```

In [ ]: # Get unique levels of the 'age' column
#new_age_levels = df_diabetes['age'].unique()

# Print each unique level in a new line
##print("Unique levels in 'age' column:")
#for level in new_age_levels:
#     print('-', level)

```

Education

The education levels have been reduced from 6 to 3. Levels 1 and 2 assigned to level 0 (Elementary); 3 and 4 to level 1 (High school), 5 and 6 to level 2 (collage)

```

In [ ]: education_mapping = {
        1: 0, # Levels 1 and 2 map to 0 (up to grade 8)
        2: 0,
        3: 1, # Levels 3 and 4 map to 1 (grade 9 to 12th)
        4: 1,
        5: 2, # Levels 5 and 6 map to 1 (collage or university)
        6: 2
    }
## Replace the aeducation levels in the DataFrame
df_diabetes['education'] = df_diabetes['education'].map(education_mapping)

## Rename the education levels to Elementary, High School and Collage
#df_diabetes['education'] = df_diabetes['education'].replace({0: 'Elementary', 1: '

## We can skip the renaming step if we prefer using numerical values instead of st
## Now, 'Education' column will have levels 'Elementary', 'High Schhol', & 'Collage

```

```

In [ ]: # Get unique levels of the 'education' column
#new_education_levels = df_diabetes['education'].unique()

# Print each unique level in a new line
#print("Unique levels in 'education' column:")
#for level in new_education_levels:
#     print('-', level)

```

Income

The income levels have been reduced from 8 to 3. Levels 1 to 5 assigned to level 0 (Low income); 6 and 7 to level 1 (Middle income), 8 to level 2 (High income)

```

In [ ]: income_mapping = {
        1: 0, # Levels 1 and 5 map to 0 (Low income)

```

```

2: 0,
3: 0,
4: 0,
5: 0,
6: 1, # Levels 6 and 7 map to 1 (middle income)
7: 1,
8: 2  # Levels 8 map to 2 (high income)
}

## Then replace the income levels in the DataFrame
df_diabetes['income'] = df_diabetes['income'].map(income_mapping)

## Finally, rename the income levels to low income, middle income, and high income

# df_diabetes['income'] = df_diabetes['income'].replace({0: 'Low income', 1: 'Middle income', 2: 'High income'})
## we can skip the renaming step if we prefer using numerical values instead of strings
## Now 'income' column will have levels 'low income', 'middle income', 'high income'

```

```

In [ ]: # Get unique levels of the 'education' column
# new_income_levels = df_diabetes['income'].unique()

# Print each unique level in a new line
#print("Unique levels in 'income' column:")
#for level in new_income_levels:
#    print('-', level)

```

6. Frequency distribution of the target/outcome variable, diabetes

```

In [ ]: import pandas as pd
class_diabetes = df_diabetes['diabetes'].value_counts()

# Display the class frequencies
print("class diabetes:")
print(class_diabetes)

```

```

class diabetes:
diabetes
1    35097
0    33960
Name: count, dtype: int64

```

```


In [ ]: df_diabetes['diabetes'].value_counts().plot(kind='bar')

```

```

Out[ ]: <Axes: xlabel='diabetes'>

```



7. Outliers: Detect, visualize and remove outliers.

7.1. Detect outliers

The lower and upper boundaries beyond which data points are considered outliers are given as follows

- The lower bound is the value at the first quartile (Q1) minus 1.5 times the interquartile range (IQR).
 - Lower Bound = $Q1 - 1.5 \times IQR$
- The upper bound is the value at the third quartile (Q3) plus 1.5 times the interquartile range (IQR).
 - Upper Bound = $Q3 + 1.5 \times IQR$

Detect outliers using IQR and count them by features

```
In [ ]: import pandas as pd

# Function to detect outliers using IQR and count them by feature
def detect_outliers_iqr(df_diabetes, columns_with_outliers):
    outliers_count = {col: 0 for col in columns_with_outliers}

    for col in columns_with_outliers:
        q1 = df_diabetes[col].quantile(0.25)
        q3 = df_diabetes[col].quantile(0.75)
        iqr = q3 - q1
        lower_bound = q1 - 1.5 * iqr
```

```

upper_bound = q3 + 1.5 * iqr
for val in df_diabetes[col]:
    if val < lower_bound or val > upper_bound:
        outliers_count[col] += 1

return outliers_count

columns_with_outliers = ['bmi', 'menthlth', 'phyhlth']
outliers_count = detect_outliers_iqr(df_diabetes, columns_with_outliers)

# Create a DataFrame to tabulate the number of outliers by feature
outliers_count_df = pd.DataFrame(list(outliers_count.items()), columns=['Feature',

print("\nNumber of outliers by feature:")
print(outliers_count_df)

```

Number of outliers by feature:

	Feature	Number of Outliers
0	bmi	2181
1	menthlth	10703
2	phyhlth	10620

7.2.Vizualize Outliers using Box Plots

```

In [ ]: # List of features to plot
columns_with_outliers = ['bmi', 'menthlth', 'phyhlth']



# Loop through each feature and create a box plot
for feature in columns_with_outliers:
    plt.figure(figsize=(10, 6))

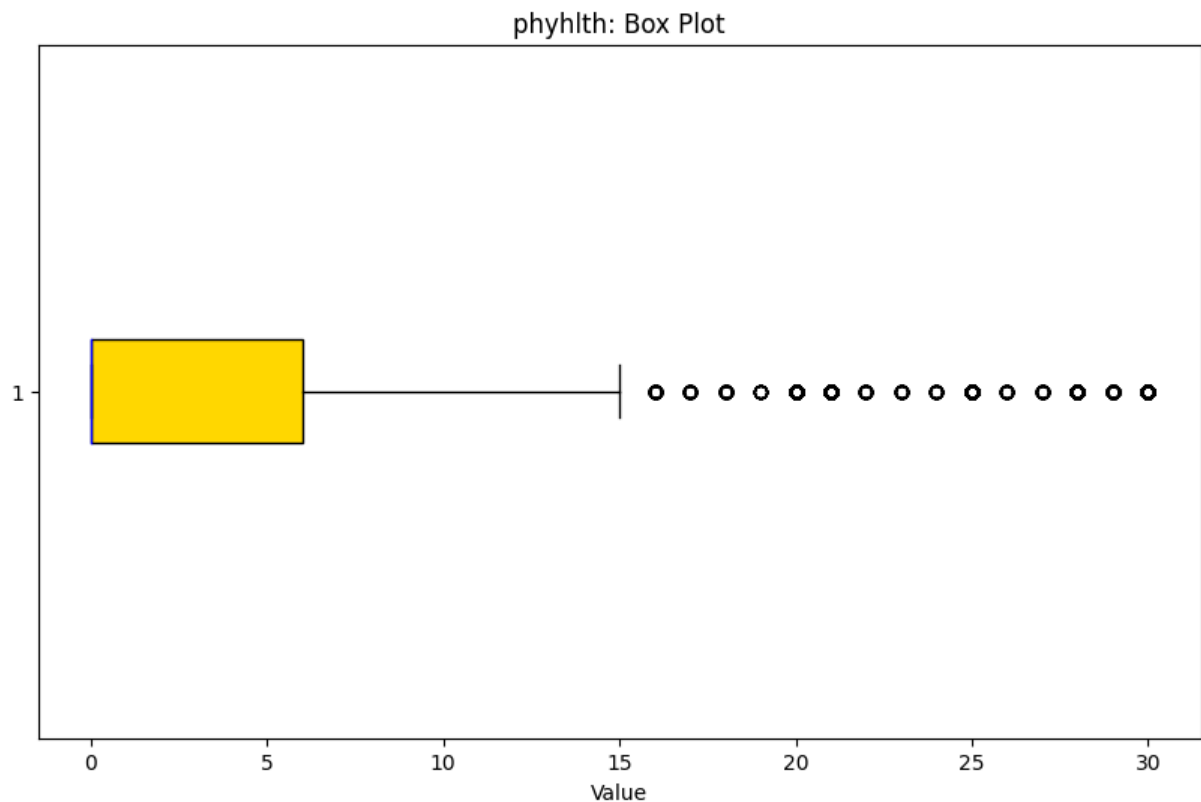
    # Create box plot
    plt.boxplot(df_diabetes[feature], vert=False, patch_artist=True, boxprops=dict(

    # Add Labels and title
    plt.xlabel('Value')
    plt.title(f'{feature}: Box Plot')

    # Show the plot
    plt.show()

```





7.3. Detect and drop outliers from the dataframe

- Drop the outliers using the indices of the outliers obtained from the `detect_outliers_iqr` function

```
In [ ]: import pandas as pd

# Function to detect outliers using IQR for specified columns
def detect_outliers_iqr(df_diabetes, columns_with_outliers):
    outliers = set()
    for col in columns_with_outliers:
        q1 = df_diabetes[col].quantile(0.25)
        q3 = df_diabetes[col].quantile(0.75)
        iqr = q3 - q1
        lower_bound = q1 - 1.5 * iqr
        upper_bound = q3 + 1.5 * iqr
        for i, val in enumerate(df_diabetes[col]):
            if val < lower_bound or val > upper_bound:
                outliers.add(i) # Add index to the set of outliers
    return list(outliers)

columns_with_outliers = ['bmi', 'menthlth', 'phyhlth']
outliers_indices = detect_outliers_iqr(df_diabetes, columns_with_outliers)

# Drop outliers from the DataFrame
#diabetes_clean = diabetes_clean.drop(outliers_indices)
#
# Print the cleaned DataFrame
```

```
#print("DataFrame after removing outliers:")
# print(diabetes_clean)
```

8.Histogram- numerical features

```
In [ ]: #import pandas as pd
#import numpy as np
#import matplotlib.pyplot as plt
#from scipy.stats import norm

# List of features to plot
features = ['bmi', 'menthlth', 'phyhlth']

# Loop through each feature and create histogram with normal distribution fit
for feature in features:
    plt.figure(figsize=(10, 6))

    # Create histogram
    plt.hist(df_diabetes[feature], bins=20, color='limegreen', edgecolor='black', a



    # Fit a normal distribution to the data
    mu, std = norm.fit(df_diabetes[feature])


    # Overlay normal distribution curve
    xmin, xmax = plt.xlim()
    x_range = np.linspace(xmin, xmax, 100)
    p = norm.pdf(x_range, mu, std)

    plt.plot(x_range, p, 'blue', linewidth=2, label=f'Fit (mu={mu:.2f}, std={std:.2

    # Add labels and title
    plt.xlabel(feature)
    plt.ylabel('Density')
    plt.title(f'{feature}: Histogram with Normal Distribution Fit')
    plt.legend()

    # Show the plot
    plt.show()
```



9. Scatter Plot/Bar Plot: used to explore relationships between variables.

```
In [ ]: # List of features to plot
features = ['bmi', 'menthlth', 'phyhlth']

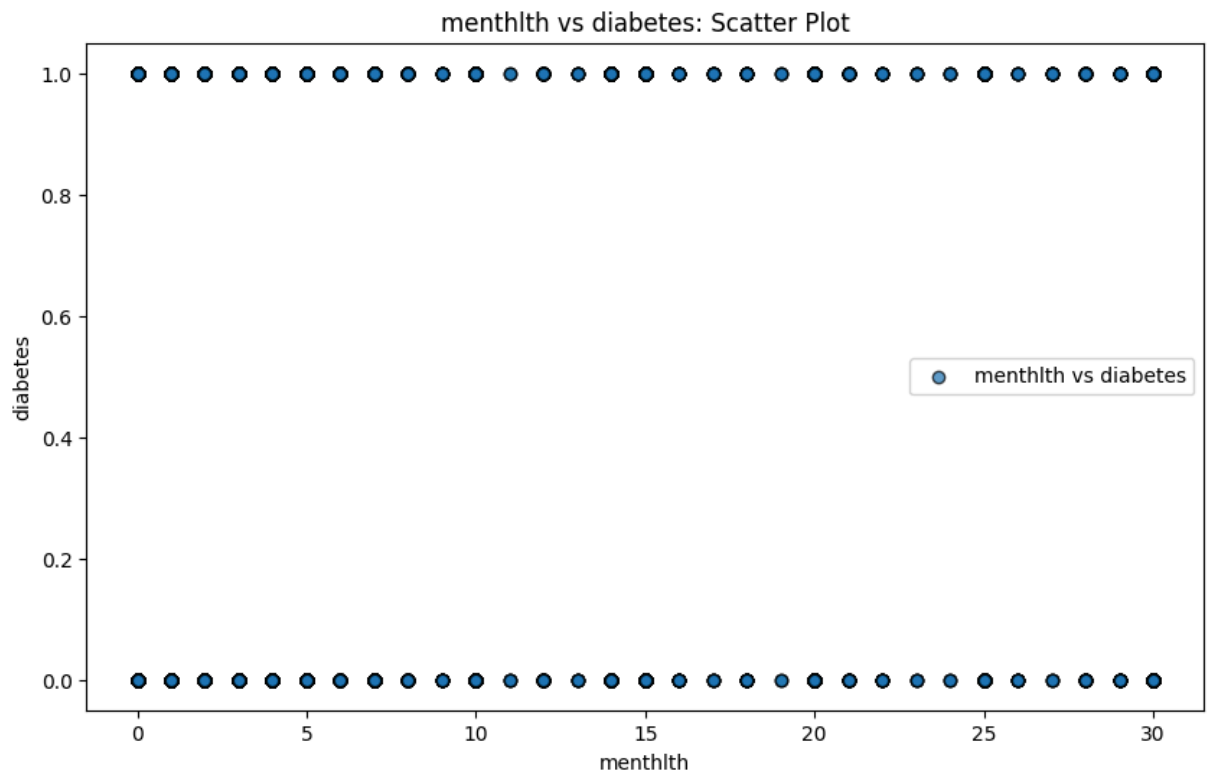
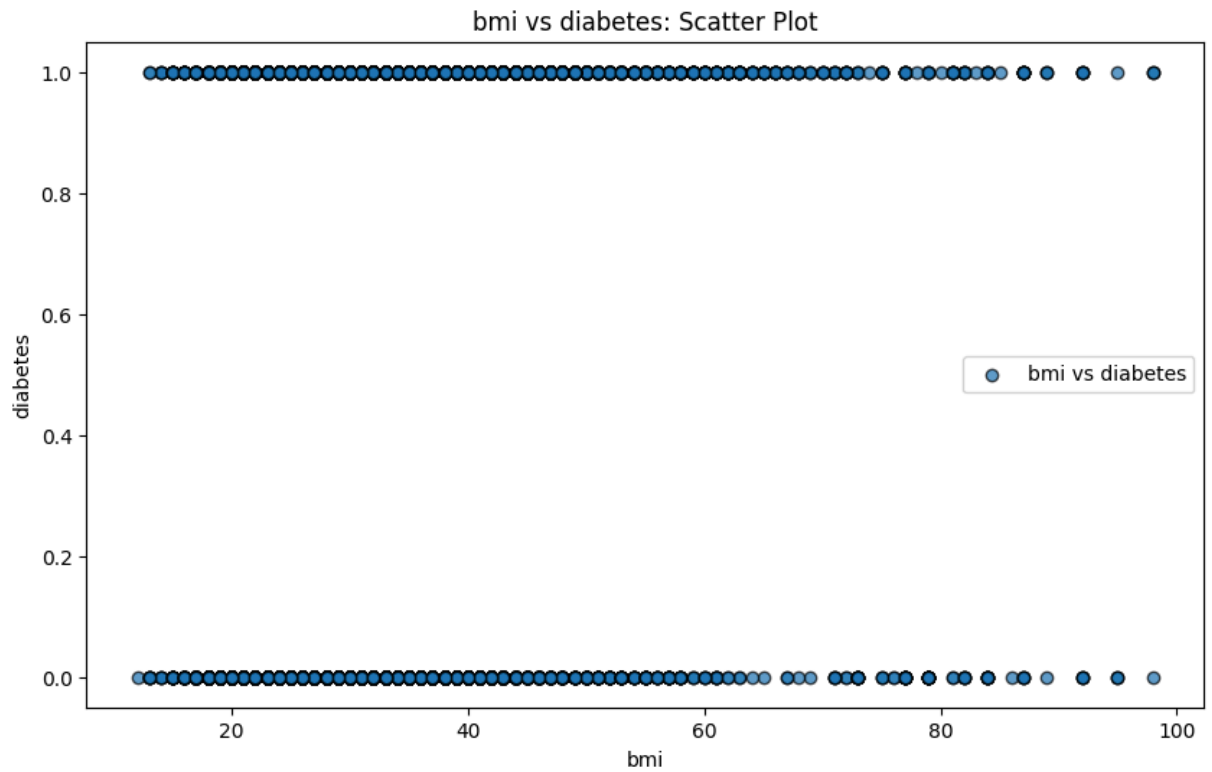
# Target variable
target = 'diabetes'

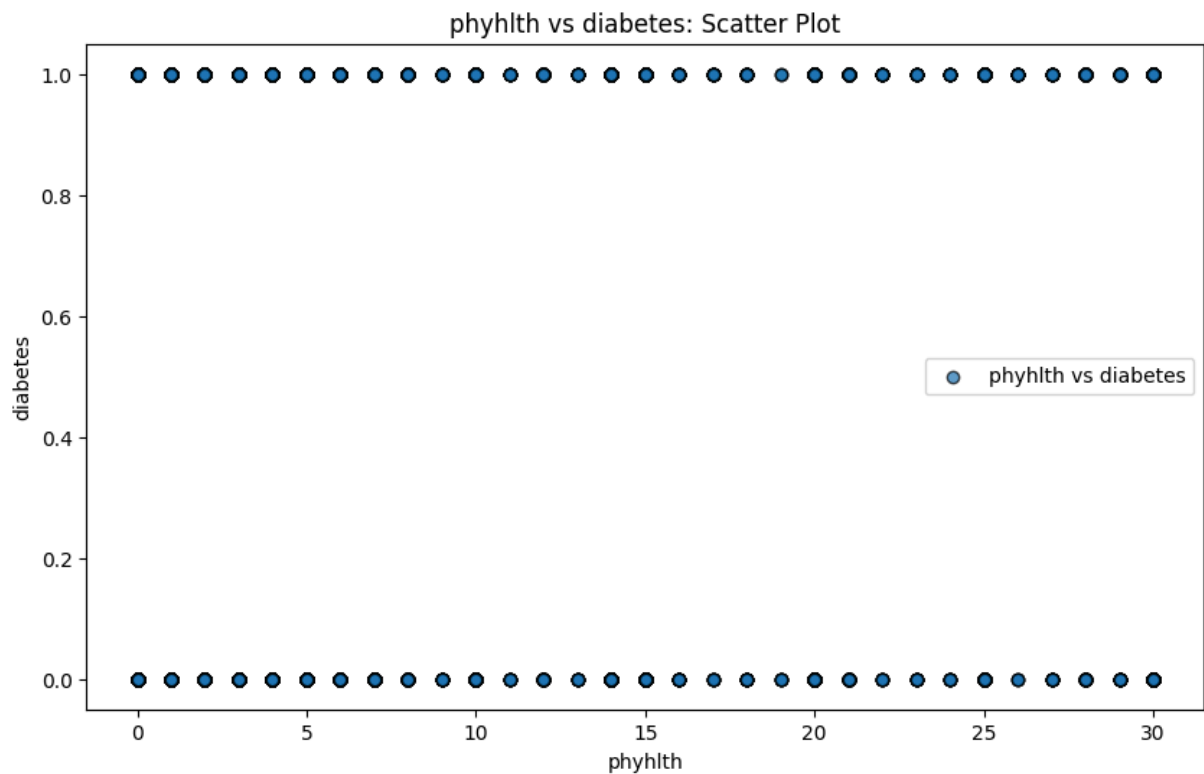
# Loop through each feature and create a scatter plot
for feature in features:
    plt.figure(figsize=(10, 6))

    # Create scatter plot
    plt.scatter(df_diabetes[feature], df_diabetes[target], alpha=0.7, edgecolor='b')

    # Add labels and title
    plt.xlabel(feature)
    plt.ylabel(target)
    plt.title(f'{feature} vs {target}: Scatter Plot')
    plt.legend()

# Show the plot
plt.show()
```





10. Summary Statistics: numeric features.

```
In [ ]: # df_diabetes.describe()

# For readability, transpose the summary statistics
df_diabetes.describe().T
```

Out[]:

	count	mean	std	min	25%	50%	75%	max
diabetes	69057.0	0.508232	0.499936	0.0	0.0	1.0	1.0	1.0
bp	69057.0	0.571224	0.494905	0.0	0.0	1.0	1.0	1.0
chol	69057.0	0.531329	0.499021	0.0	0.0	1.0	1.0	1.0
cholcheck	69057.0	0.974803	0.156723	0.0	1.0	1.0	1.0	1.0
bmi	69057.0	29.955834	7.147972	12.0	25.0	29.0	33.0	98.0
smoker	69057.0	0.481935	0.499677	0.0	0.0	0.0	1.0	1.0
stroke	69057.0	0.063643	0.244118	0.0	0.0	0.0	0.0	1.0
heart_disease	69057.0	0.150875	0.357930	0.0	0.0	0.0	0.0	1.0
activity	69057.0	0.696483	0.459780	0.0	0.0	1.0	1.0	1.0
fruits	69057.0	0.605659	0.488712	0.0	0.0	1.0	1.0	1.0
veggies	69057.0	0.784120	0.411434	0.0	1.0	1.0	1.0	1.0
alcohol	69057.0	0.043515	0.204014	0.0	0.0	0.0	0.0	1.0
healthcare	69057.0	0.953908	0.209687	0.0	1.0	1.0	1.0	1.0
nodocbccost	69057.0	0.096138	0.294782	0.0	0.0	0.0	0.0	1.0
genhlth	69057.0	2.863692	1.107950	1.0	2.0	3.0	4.0	5.0
menthlth	69057.0	3.840103	8.231164	0.0	0.0	0.0	3.0	30.0
phyhlth	69057.0	5.945306	10.139113	0.0	0.0	0.0	6.0	30.0
walk	69057.0	0.258612	0.437875	0.0	0.0	0.0	1.0	1.0
sex	69057.0	0.456464	0.498105	0.0	0.0	0.0	1.0	1.0
age	69057.0	2.174160	0.843393	0.0	2.0	2.0	3.0	3.0
education	69057.0	1.619329	0.534450	0.0	1.0	2.0	2.0	2.0
income	69057.0	0.870049	0.819899	0.0	0.0	1.0	2.0	2.0

11. Correlation Matrix:

Visualize the correlation between features & between features and the outcome variable(diabetes)

```
In [ ]: ## Compute the correlation matrix including the outcome variable, 'diabetes'
correlation_matrix = df_diabetes.corr()

## Compute the correlation matrix by dropping the outcome variable, 'diabetes'
# df_diabetes_drop_target = df_diabetes.drop(columns=['diabetes'])
# correlation_matrix = df_diabetes_drop_target.corr()
```

```
# Display the correlation matrix  
print(correlation_matrix)
```

	diabetes	bp	chol	cholcheck	bmi	smoker	\
diabetes	1.000000	0.372048	0.281399	0.118900	0.285643	0.075853	
bp	0.372048	1.000000	0.308987	0.106593	0.232372	0.078123	
chol	0.281399	0.308987	1.000000	0.088231	0.123917	0.086522	
cholcheck	0.118900	0.106593	0.088231	1.000000	0.047779	-0.002854	
bmi	0.285643	0.232372	0.123917	0.047779	1.000000	0.002761	
smoker	0.075853	0.078123	0.086522	-0.002854	0.002761	1.000000	
stroke	0.122727	0.126869	0.098166	0.023368	0.019503	0.061957	
heart_disease	0.207229	0.206776	0.178207	0.044795	0.055345	0.120457	
activity	-0.150281	-0.128307	-0.084469	-0.010072	-0.164179	-0.072401	
fruits	-0.044560	-0.031818	-0.040783	0.015853	-0.076933	-0.068192	
veggies	-0.072181	-0.059824	-0.037801	-0.001040	-0.050163	-0.023760	
alcohol	-0.098709	-0.029764	-0.027259	-0.026850	-0.060795	0.076394	
healthcare	0.027034	0.039659	0.034352	0.106549	-0.010527	-0.010228	
nodocbccost	0.036145	0.021802	0.029976	-0.061975	0.061861	0.031896	
genhlth	0.396571	0.308459	0.227588	0.063116	0.256642	0.140658	
menthlth	0.080688	0.058133	0.079929	-0.009365	0.099286	0.086354	
phyhlth	0.206868	0.167821	0.138266	0.036442	0.155661	0.114730	
walk	0.267082	0.229638	0.157859	0.046421	0.240667	0.113713	
sex	0.042538	0.037824	0.013250	-0.008116	-0.002822	0.113422	
age	0.270581	0.325906	0.236471	0.100686	-0.026132	0.108439	
education	-0.131569	-0.109807	-0.067133	-0.012338	-0.061120	-0.094346	
income	-0.210774	-0.177433	-0.095025	0.003931	-0.105602	-0.098975	

	stroke	heart_disease	activity	fruits	...	healthcare	\
diabetes	0.122727	0.207229	-0.150281	-0.044560	...	0.027034	
bp	0.126869	0.206776	-0.128307	-0.031818	...	0.039659	
chol	0.098166	0.178207	-0.084469	-0.040783	...	0.034352	
cholcheck	0.023368	0.044795	-0.010072	0.015853	...	0.106549	
bmi	0.019503	0.055345	-0.164179	-0.076933	...	-0.010527	
smoker	0.061957	0.120457	-0.072401	-0.068192	...	-0.010228	
stroke	1.000000	0.222062	-0.076771	-0.005811	...	0.007801	
heart_disease	0.222062	1.000000	-0.093858	-0.014931	...	0.017603	
activity	-0.076771	-0.093858	1.000000	0.127578	...	0.024168	
fruits	-0.005811	-0.014931	0.127578	1.000000	...	0.026964	
veggies	-0.044869	-0.032327	0.143392	0.234505	...	0.026832	
alcohol	-0.024496	-0.038745	0.021624	-0.031518	...	-0.012691	
healthcare	0.007801	0.017603	0.024168	0.026964	...	1.000000	
nodocbccost	0.034305	0.033397	-0.059079	-0.042215	...	-0.220451	
genhlth	0.186537	0.271502	-0.264142	-0.086836	...	-0.028477	
menthlth	0.084800	0.071530	-0.124535	-0.057019	...	-0.047689	
phyhlth	0.161824	0.194963	-0.228329	-0.041836	...	-0.000416	
walk	0.189714	0.229188	-0.270988	-0.044017	...	0.011066	
sex	0.004149	0.099020	0.052069	-0.088017	...	-0.006804	
age	0.113841	0.205764	-0.088975	0.055290	...	0.144226	
education	-0.056268	-0.075456	0.154878	0.075876	...	0.090365	
income	-0.120743	-0.134466	0.182744	0.056529	...	0.117616	

	nodocbccost	genhlth	menthlth	phyhlth	walk	sex	\
diabetes	0.036145	0.396571	0.080688	0.206868	0.267082	0.042538	
bp	0.021802	0.308459	0.058133	0.167821	0.229638	0.037824	
chol	0.029976	0.227588	0.079929	0.138266	0.157859	0.013250	
cholcheck	-0.061975	0.063116	-0.009365	0.036442	0.046421	-0.008116	
bmi	0.061861	0.256642	0.099286	0.155661	0.240667	-0.002822	
smoker	0.031896	0.140658	0.086354	0.114730	0.113713	0.113422	
stroke	0.034305	0.186537	0.084800	0.161824	0.189714	0.004149	

heart_disease	0.033397	0.271502	0.071530	0.194963	0.229188	0.099020
activity	-0.059079	-0.264142	-0.124535	-0.228329	-0.270988	0.052069
fruits	-0.042215	-0.086836	-0.057019	-0.041836	-0.044017	-0.088017
veggies	-0.033643	-0.106136	-0.047458	-0.060875	-0.078098	-0.053422
alcohol	0.008453	-0.063705	0.013914	-0.038739	-0.051894	0.015437
healthcare	-0.220451	-0.028477	-0.047689	-0.000416	0.011066	-0.006804
nodocbccost	1.000000	0.164758	0.191108	0.153952	0.123415	-0.048469
genhlth	0.164758	1.000000	0.310093	0.550138	0.472338	-0.016880
menthlth	0.191108	0.310093	1.000000	0.376625	0.246948	-0.089926
phyhlth	0.153952	0.550138	0.376625	1.000000	0.484092	-0.045929
walk	0.123415	0.472338	0.246948	0.484092	1.000000	-0.082858
sex	-0.048469	-0.016880	-0.089926	-0.045929	-0.082858	1.000000
age	-0.131942	0.137805	-0.102646	0.075041	0.176205	0.000452
education	-0.066338	-0.228809	-0.072687	-0.119683	-0.155786	0.021732
income	-0.176346	-0.339565	-0.168691	-0.232727	-0.298935	0.149290

	age	education	income
diabetes	0.270581	-0.131569	-0.210774
bp	0.325906	-0.109807	-0.177433
chol	0.236471	-0.067133	-0.095025
cholcheck	0.100686	-0.012338	0.003931
bmi	-0.026132	-0.061120	-0.105602
smoker	0.108439	-0.094346	-0.098975
stroke	0.113841	-0.056268	-0.120743
heart_disease	0.205764	-0.075456	-0.134466
activity	-0.088975	0.154878	0.182744
fruits	0.055290	0.075876	0.056529
veggies	-0.015007	0.133288	0.138043
alcohol	-0.052377	0.035449	0.070686
healthcare	0.144226	0.090365	0.117616
nodocbccost	-0.131942	-0.066338	-0.176346
genhlth	0.137805	-0.228809	-0.339565
menthlth	-0.102646	-0.072687	-0.168691
phyhlth	0.075041	-0.119683	-0.232727
walk	0.176205	-0.155786	-0.298935
sex	0.000452	0.021732	0.149290
age	1.000000	-0.087148	-0.149126
education	-0.087148	1.000000	0.355316
income	-0.149126	0.355316	1.000000

[22 rows x 22 columns]

Correlation Heatmap


```
In [ ]: #import pandas as pd
        #import matplotlib.pyplot as plt
        #import seaborn as sns

        # Run correlation matrix and plot
        f, ax = plt.subplots(figsize=(10, 8))
        corr = df_diabetes.corr()

        sns.heatmap(corr, mask=np.zeros_like(corr, dtype=bool),
                    cmap=sns.diverging_palette(220, 10, as_cmap=True),
                    square=True, ax=ax)
```



```
plt.show()
```



Correlation of features with the outcome variable, diabetes

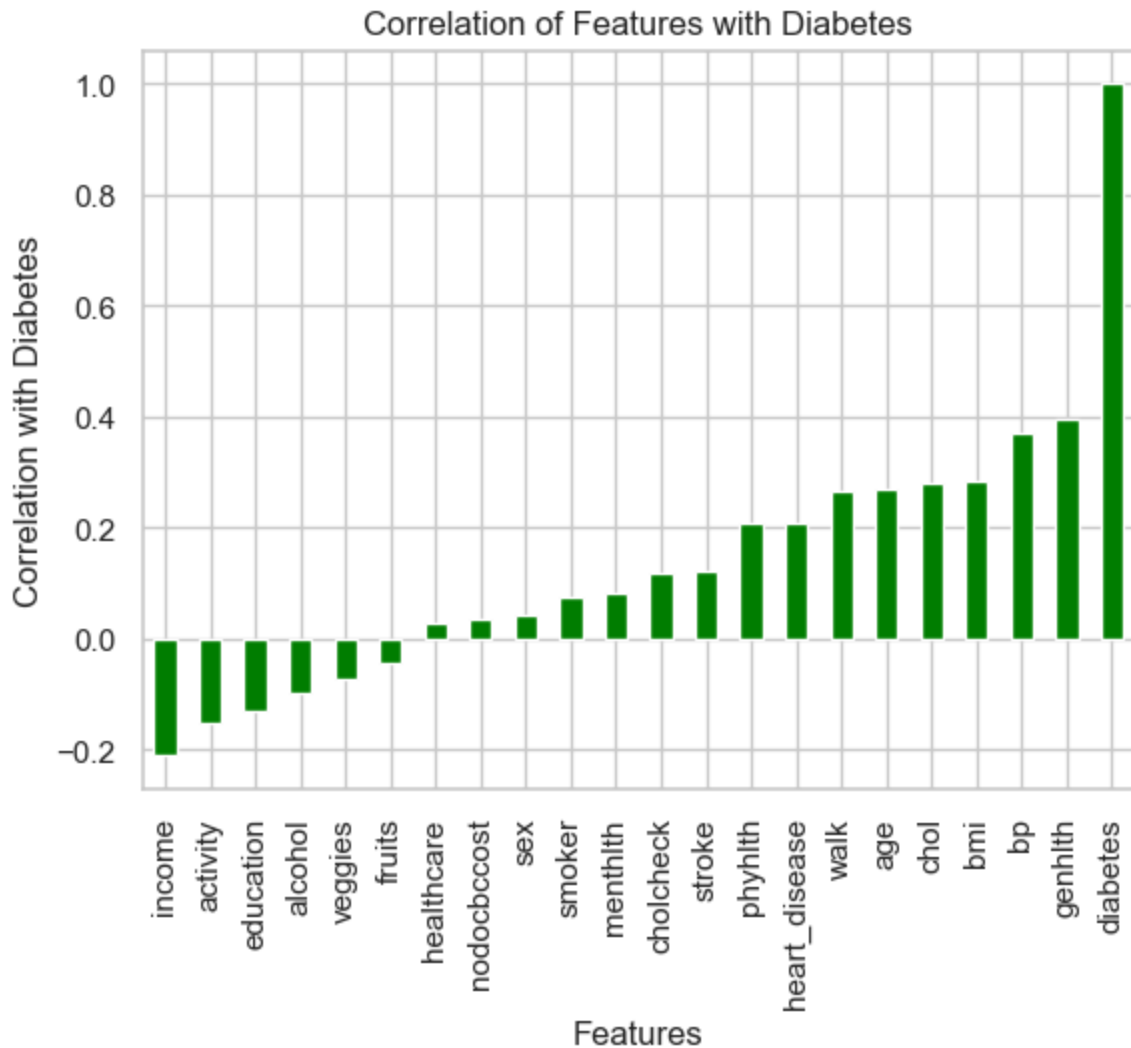
- sort the results, ascending order

```
In [ ]: df_diabetes.corr()['diabetes'].sort_values()
```

```
Out[ ]: income      -0.210774
        activity    -0.150281
        education   -0.131569
        alcohol     -0.098709
        veggies     -0.072181
        fruits      -0.044560
        healthcare  0.027034
        nodocbccost 0.036145
        sex         0.042538
        smoker      0.075853
        menthlth    0.080688
        cholcheck   0.118900
        stroke      0.122727
        phyhlth     0.206868
        heart_disease 0.207229
        walk        0.267082
        age         0.270581
        chol        0.281399
        bmi         0.285643
        bp          0.372048
        genhlth     0.396571
        diabetes    1.000000
        Name: diabetes, dtype: float64
```

Bar graph: Correlation of features with the outcome variable, Diabetes

```
In [ ]: df_diabetes.corr()['diabetes'].sort_values().plot(kind='bar', color='green')
        plt.xlabel('Features')
        plt.ylabel('Correlation with Diabetes')
        plt.title('Correlation of Features with Diabetes')
        plt.show()
```



12. Chi-square test for independence of categorical variables against 'diabetes'

- The null hypothesis (H_0) in Chi-square test of independence states that there is no relationship between the categorical variables being tested.
- Given a significance level of 0.05, if we reject the null hypothesis(False), it indicates that there is evidence to suggest a relationship between the variables.

```
In [ ]: import pandas as pd
from scipy.stats import chi2_contingency

# List to store the results
chi2_results = []

# Define the significance level
alpha = 0.05

df_diabetes_drop_target = df_diabetes.drop(columns=['diabetes'])

# Iterate over each column in the DataFrame
for column in df_diabetes.columns:
```

```

# Create a contingency table between the current column and the target column
contingency_table = pd.crosstab(df_diabetes[column], df_diabetes['diabetes'])

# Perform the Chi-square test
chi2_stat, p_val, _, _ = chi2_contingency(contingency_table)

# Store the results
chi2_results.append((column, chi2_stat, p_val, p_val > alpha))

# Convert the results to a DataFrame for easier analysis
chi2_results_df = pd.DataFrame(chi2_results, columns=['Column', 'Chi-square statistic', 'p-value', 'Null Hypothesis (Ho)'])

# Print the DataFrame
print(" Chi-square Test of Independence")

print(chi2_results_df)

```

Chi-square Test of Independence

	Column	Chi-square statistic	p-value	Null Hypothesis (Ho)
0	diabetes	69052.998973	0.000000e+00	False
1	bp	9557.323086	0.000000e+00	False
2	chol	5467.178055	0.000000e+00	False
3	cholcheck	974.756080	5.513841e-214	False
4	bmi	7206.276100	0.000000e+00	False
5	smoker	397.032404	2.437524e-88	False
6	stroke	1039.121611	5.633771e-228	False
7	heart_disease	2964.403557	0.000000e+00	False
8	activity	1558.959632	0.000000e+00	False
9	fruits	136.935677	1.245551e-31	False
10	veggies	359.443376	3.722096e-80	False
11	alcohol	671.888634	3.880471e-148	False
12	healthcare	50.212286	1.379812e-12	False
13	nodocbccost	89.975138	2.411716e-21	False
14	genhlth	11355.569296	0.000000e+00	False
15	menthlth	596.821393	1.347932e-106	False
16	phyhlth	3362.803539	0.000000e+00	False
17	walk	4924.796104	0.000000e+00	False
18	sex	124.787363	5.665131e-29	False
19	age	5345.689918	0.000000e+00	False
20	education	1195.820177	2.142633e-260	False
21	income	3069.107535	0.000000e+00	False

13. Frequency distribution of some categorical variables: 'age', 'education', 'income'

```

In [ ]: import matplotlib.pyplot as plt
import seaborn as sns

# Plot the frequency distribution of each categorical variable
# categorical_columns = df_diabetes.select_dtypes(include=['category', 'object']).columns
age_education_income_features = ['age', 'education', 'income']



for column in age_education_income_features:
    plt.figure(figsize=(8, 4))
    ax = sns.countplot(data=df_diabetes, x=column, order=df_diabetes[column].value_counts().index)
    plt.title(f'Frequency Distribution of {column}')

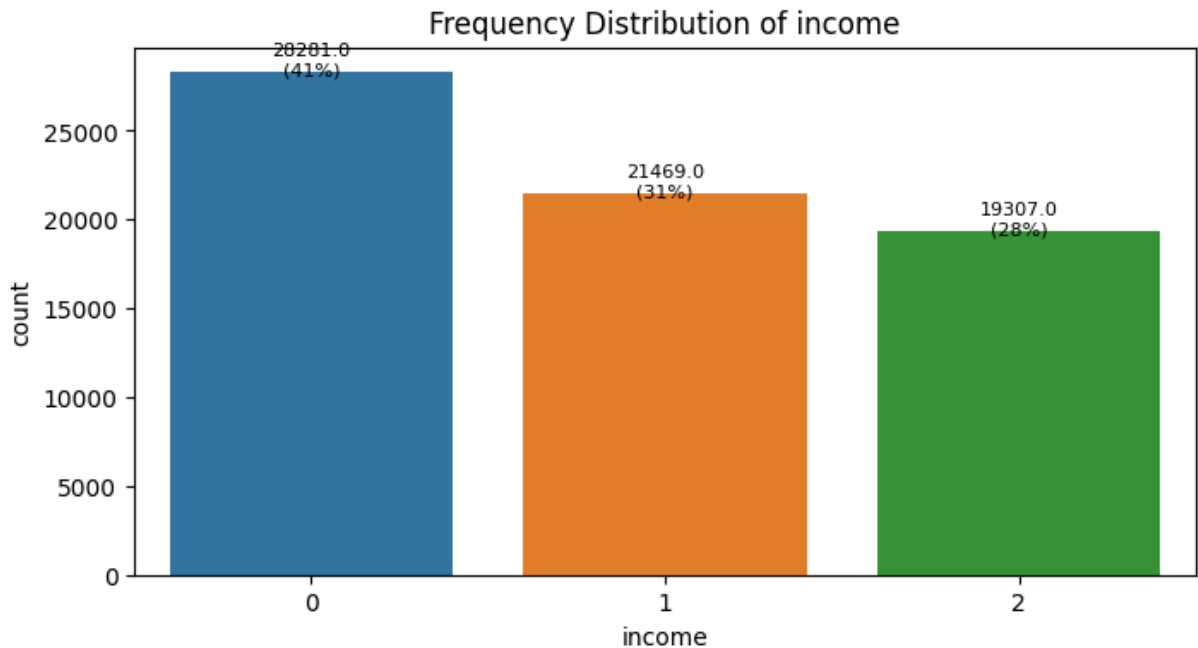
```

```
plt.xticks(rotation=0)

# Add count and percentage above each bar
total_count = len(df_diabetes[column])
for p in ax.patches:
    count = p.get_height()
    percentage = (count / total_count) * 100
    ax.annotate(f'{count}\n({percentage:.0f}%)', (p.get_x() + p.get_width() / 2,
        ha='center', va='center', fontsize=8, xytext=(0, 5), textcoords=

plt.show()
```





14.The distribution of 'diabetes' outcome by risk factors

Cross Tabulation: diabetes with bp,chol,smoker,stroke, heart_disease and alcohol

```
In [ ]: # List of features to plot
risk_facors_features = ['bp', 'chol', 'smoker', 'stroke', 'heart_disease', 'alcohol']

# Iterate through each feature and create a bar plot
for feature in risk_facors_features:
    # Create cross-tabulation
    cross_tab = pd.crosstab(df_diabetes[feature], df_diabetes['diabetes'])

    # Sum of diabetes cases for each category
    grouped = cross_tab[1]

    plt.figure(figsize=(10, 6))

    # Create alternating colors
    colors = ['tomato' if i % 2 == 0 else 'limegreen' for i in range(len(grouped))]


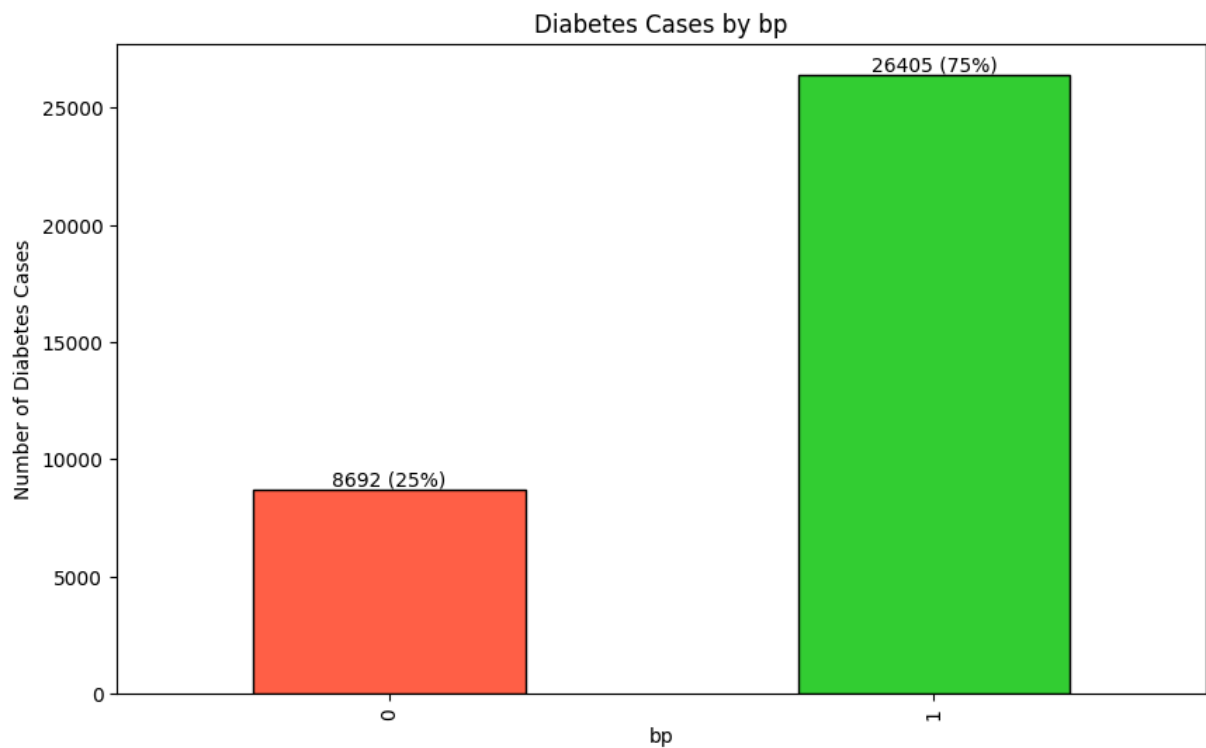
    # Plot with alternating colors
    ax = grouped.plot(kind='bar', color=colors, edgecolor='black')

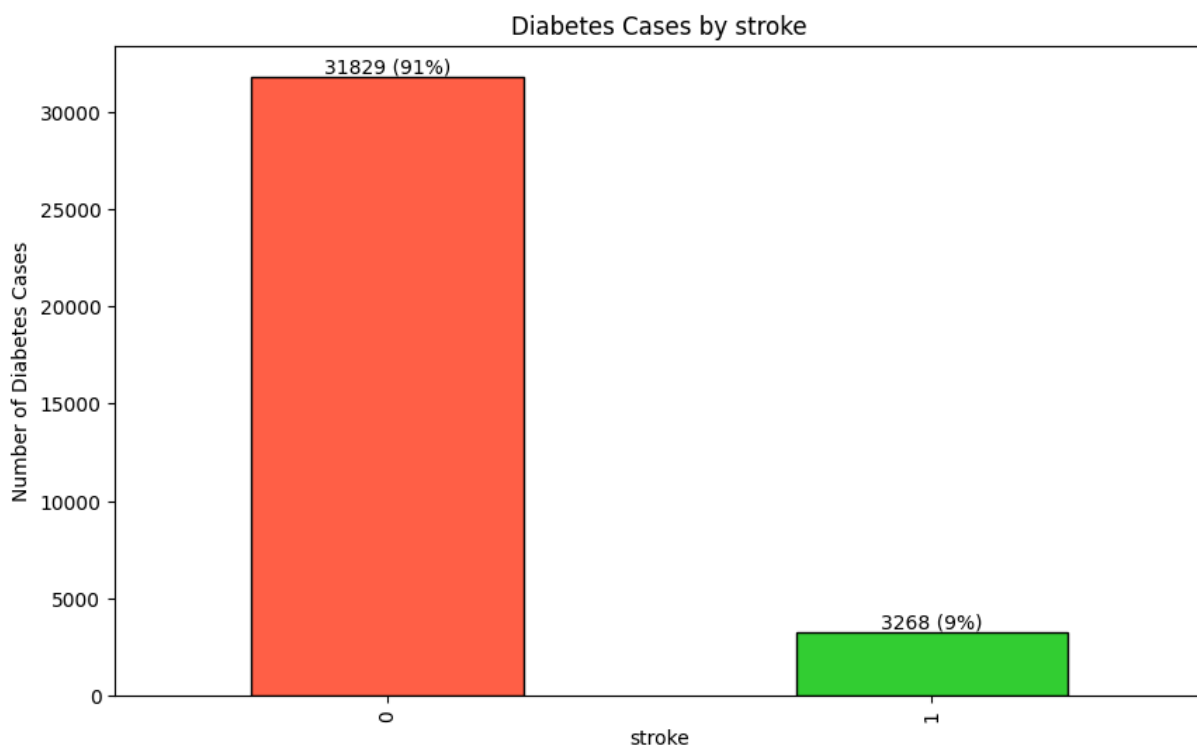
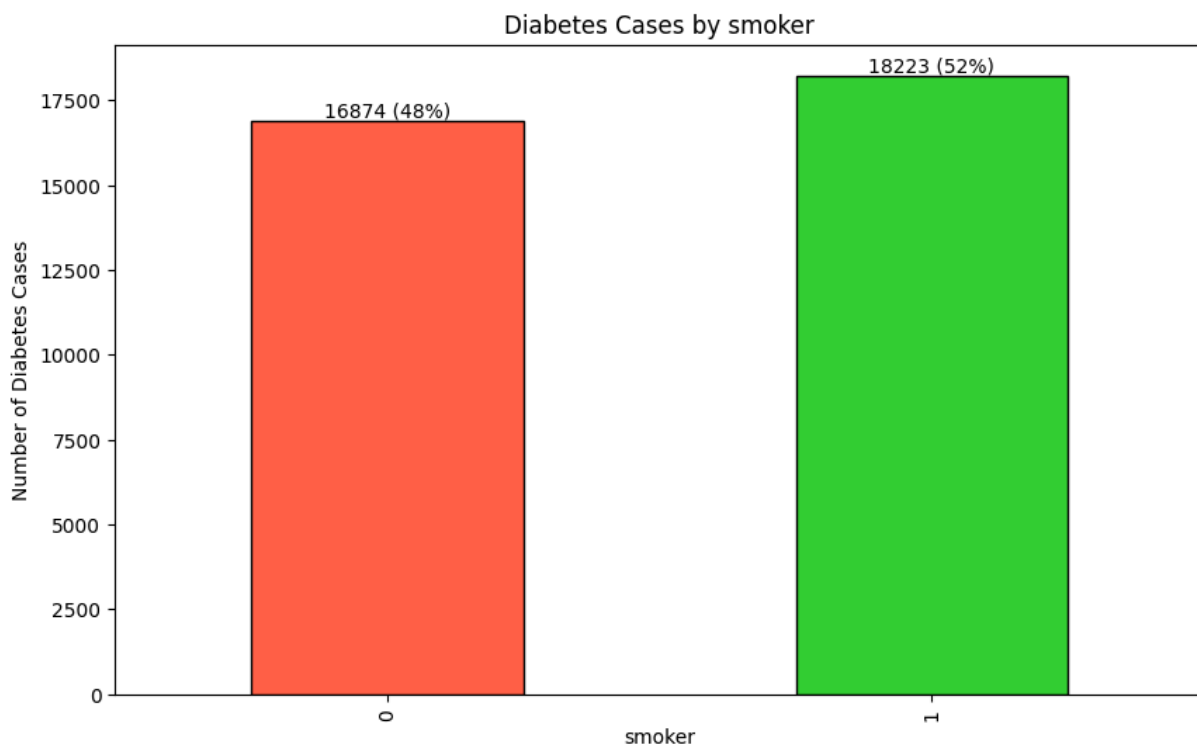
    # Calculate the total sum for percentage calculation
    total = grouped.sum()

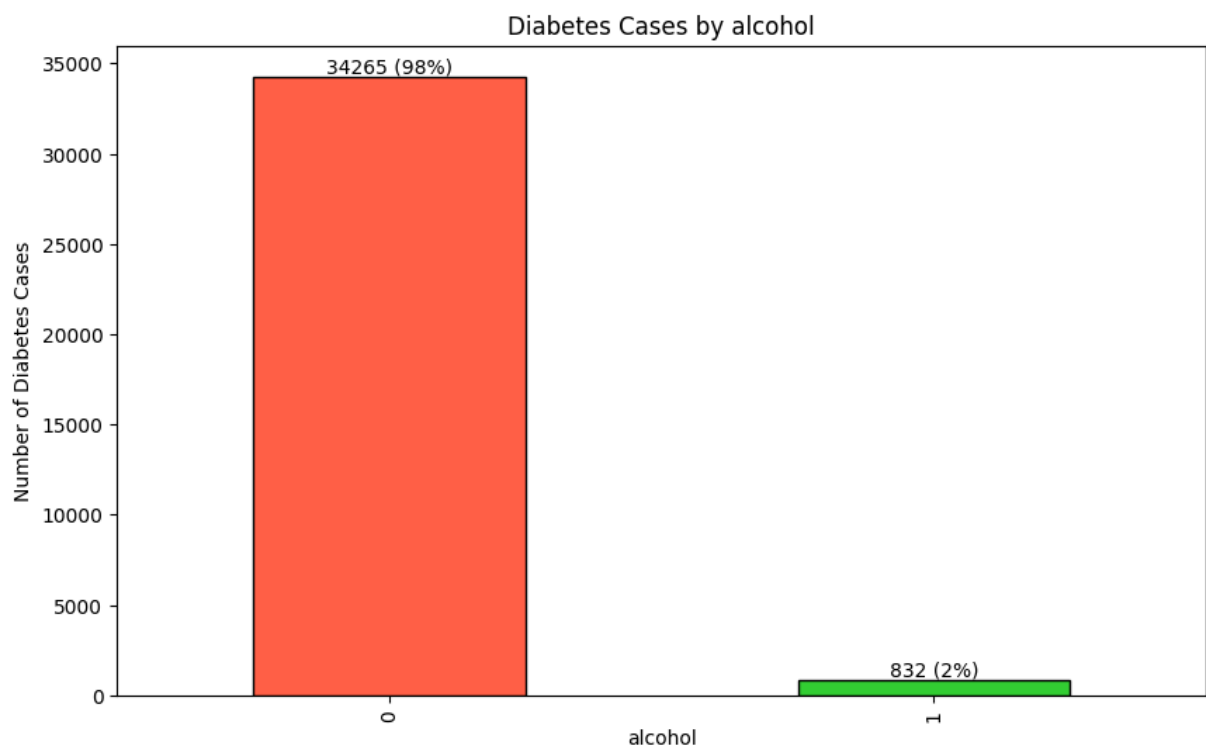
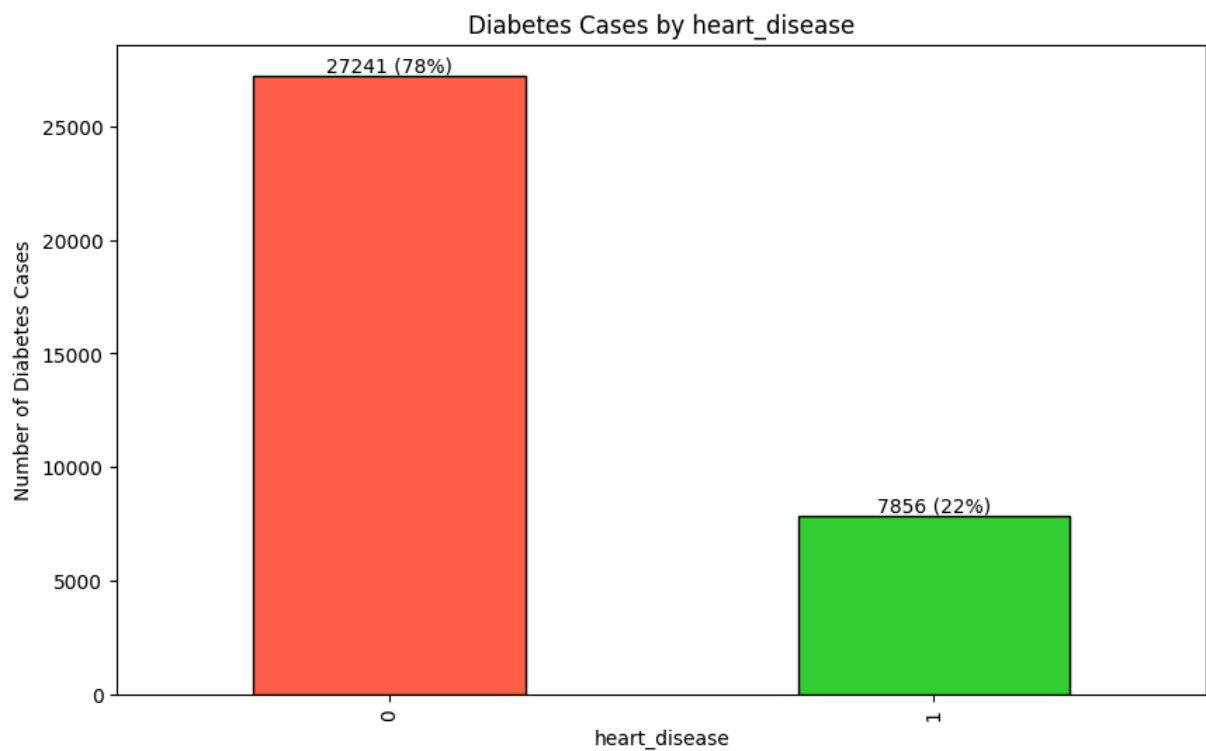
    # Annotate the bars with the sum and percentage
    for i, value in enumerate(grouped):
        percentage = (value / total) * 100
        ax.text(i, value, f'{int(round(value))} ({round(percentage)}%)', ha='center')

    # Add Labels and title
    plt.xlabel(feature)
    plt.ylabel('Number of Diabetes Cases')
    plt.title(f'Diabetes Cases by {feature}')
```

```
# Show the plot  
plt.show()
```







15. The distribution of diabetes outcome by good habit features

Croos Tabulation: diabetes with activity, fruits and veggies

```
In [ ]: # croos tabulation of good_habit_features with diabetes outcome

# List of features to generate crosstabs for
good_habit_features = ['activity', 'fruits', 'veggies']

# Iterate through each feature and create a crosstab
```

```
for feature in good_habit_features:
    crosstab = pd.crosstab(df_diabetes[feature], df_diabetes['diabetes'])
    crosstab.columns = ['No Diabetes', 'Diabetes'] # Rename columns for clarity
    print(f'Crosstab for {feature}:\n', crosstab, '\n')
```

Crosstab for activity:

	No Diabetes	Diabetes
activity		
0	7922	13038
1	26038	22059

Crosstab for fruits:

	No Diabetes	Diabetes
fruits		
0	12640	14592
1	21320	20505

Crosstab for veggies:

	No Diabetes	Diabetes
veggies		
0	6306	8602
1	27654	26495

Bar graph: Good habit features with diabetes outcome

```
In [ ]: # List of features to plot
good_habit_features = ['activity', 'fruits', 'veggies']

# Iterate through each feature and create a frequency table
for feature in good_habit_features:
    # Create a frequency table using crosstab
    frequency_table = pd.crosstab(df_diabetes[feature], df_diabetes['diabetes'])

    plt.figure(figsize=(8, 4))

    # Plot the frequency table as a bar plot
    ax = frequency_table.plot(kind='bar', stacked=True, color=['skyblue', 'tomato'])

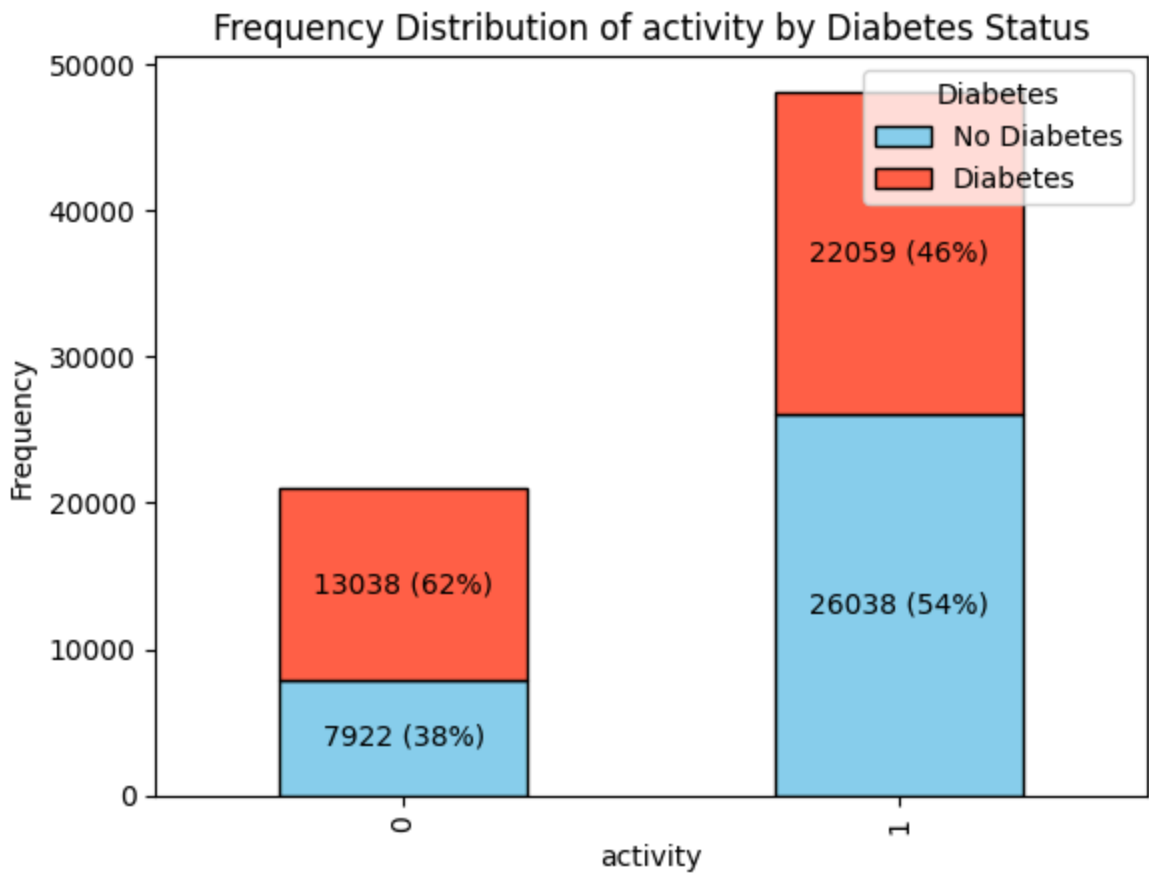
    # Annotate the bars with the counts and percentages
    for i, row in enumerate(frequency_table.values):
        for j, value in enumerate(row):
            percentage = (value / row.sum()) * 100
            ax.text(i, row[:j].sum() + value/2, f'{value} ({int(round(percentage))}%)')

    # Add labels and title
    plt.xlabel(feature)
    plt.ylabel('Frequency')
    plt.title(f'Frequency Distribution of {feature} by Diabetes Status')

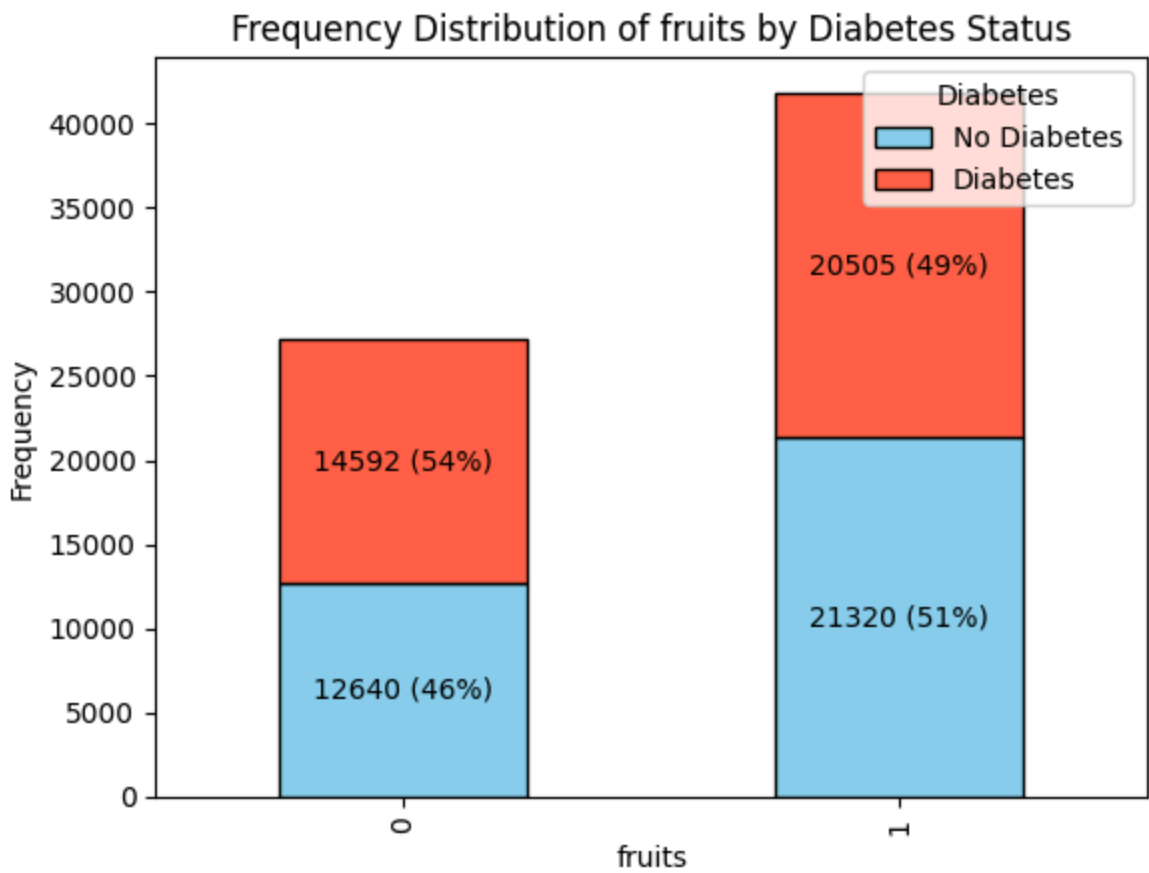
    # Add Legend
    plt.legend(title='Diabetes', labels=['No Diabetes', 'Diabetes'], loc='upper right')

    # Show the plot
    plt.show()
```


<Figure size 800x400 with 0 Axes>



<Figure size 800x400 with 0 Axes>



<Figure size 800x400 with 0 Axes>



16. Prevalence of Diabetes by age, gender, education and income level

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Plot settings
sns.set(style="whitegrid")

# List of categorical features to plot
features = ['age', 'sex', 'education', 'income']


# Creating plots using a for loop
for feature in features:
    plt.figure(figsize=(7, 4))
    prevalence = df_diabetes.groupby(feature)['diabetes'].mean().reset_index()
    prevalence['diabetes'] = prevalence['diabetes'] * 100 # Convert to percentage
    prevalence = prevalence.sort_values(by='diabetes', ascending=False) # Sort by
    sns.barplot(data=prevalence, x=feature, y='diabetes', ci=None)
    plt.title(f'prevalence of diabetes by {feature}')
    plt.ylabel('diabetes (%)')
    plt.xlabel(feature)
    plt.xticks(rotation=0)

    # Adding percentages on top of the bars
    for index, row in prevalence.iterrows():
        plt.text(index, row['diabetes'], f"{row['diabetes']:.2f}%", color='black',
```

```
plt.show()
```

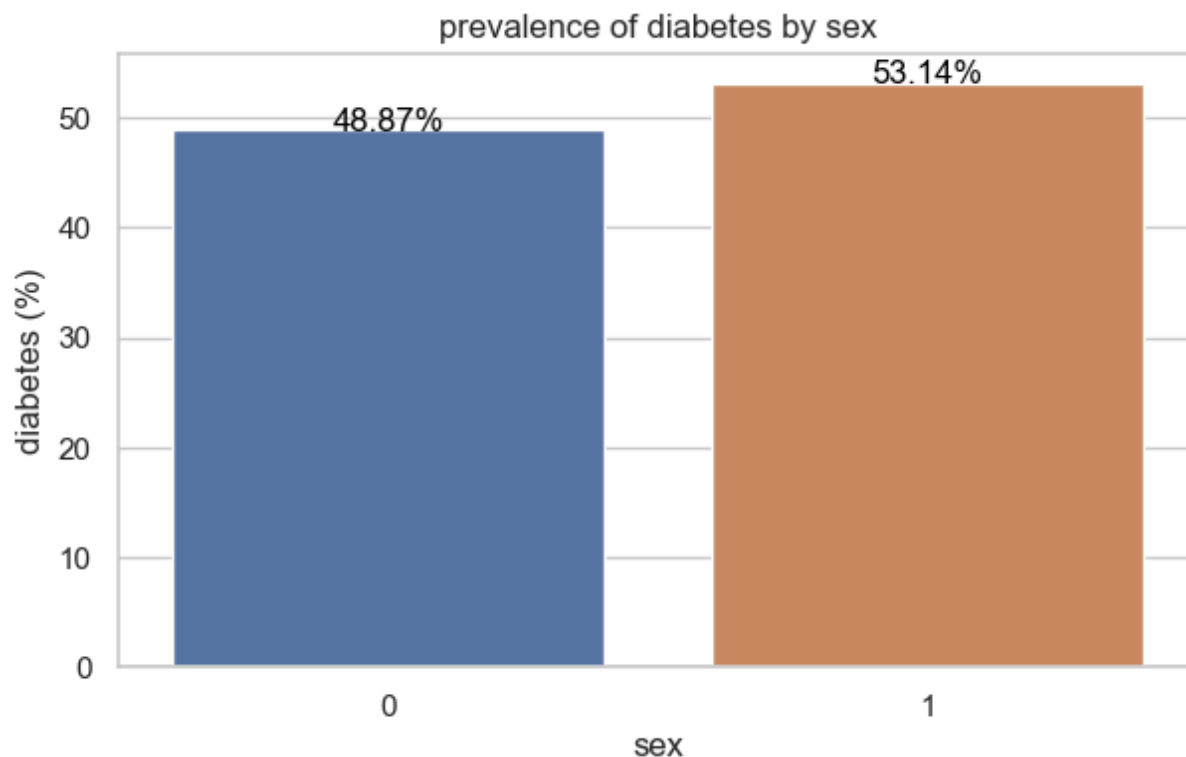
C:\Users\yitay\AppData\Local\Temp\ipykernel_8224\3554196671.py:17: FutureWarning:
The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(data=prevalence, x=feature, y='diabetes', ci=None)
```



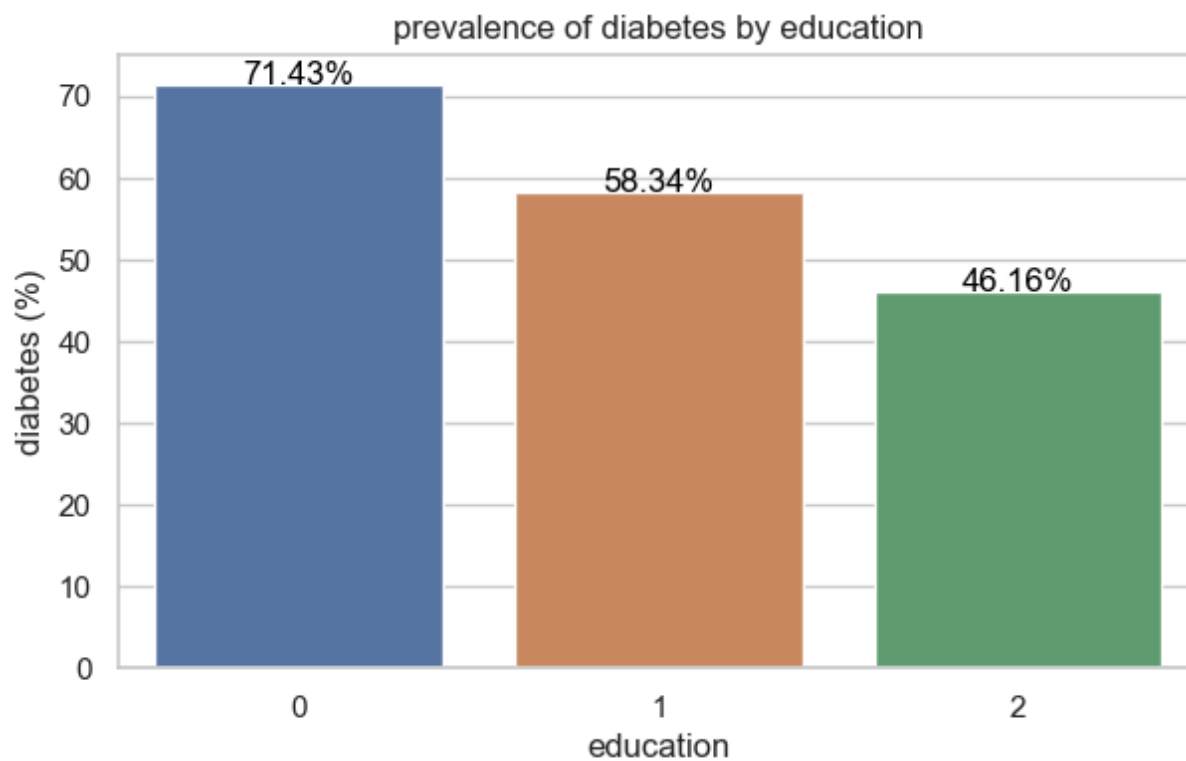
C:\Users\yitay\AppData\Local\Temp\ipykernel_8224\3554196671.py:17: FutureWarning:
The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(data=prevalence, x=feature, y='diabetes', ci=None)
```




C:\Users\yitay\AppData\Local\Temp\ipykernel_8224\3554196671.py:17: FutureWarning:
The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(data=prevalence, x=feature, y='diabetes', ci=None)
```



```
C:\Users\yitay\AppData\Local\Temp\ipykernel_8224\3554196671.py:17: FutureWarning:
The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

sns.barplot(data=prevalence, x=feature, y='diabetes', ci=None)
```



option 2

Prevalence of Diabetes by age, gender, education and income level

```
In [ ]: # List of features to plot
prevalence_features = ['age', 'sex', 'education', 'income']

# Iterate through each feature and create a frequency table
for feature in prevalence_features:
    # Create a frequency table using crosstab
    frequency_table = pd.crosstab(df_diabetes[feature], df_diabetes['diabetes'])

    plt.figure(figsize=(8, 4))

    # Plot the frequency table as a bar plot
    ax = frequency_table.plot(kind='bar', stacked=True, color=['limegreen', 'tomato'])

    # Annotate the bars with the counts and percentages
    for i, row in enumerate(frequency_table.values):
        for j, value in enumerate(row):
            percentage = (value / row.sum()) * 100
            ax.text(i, row[:j].sum() + value/2, f'{value} ({int(round(percentage))})')

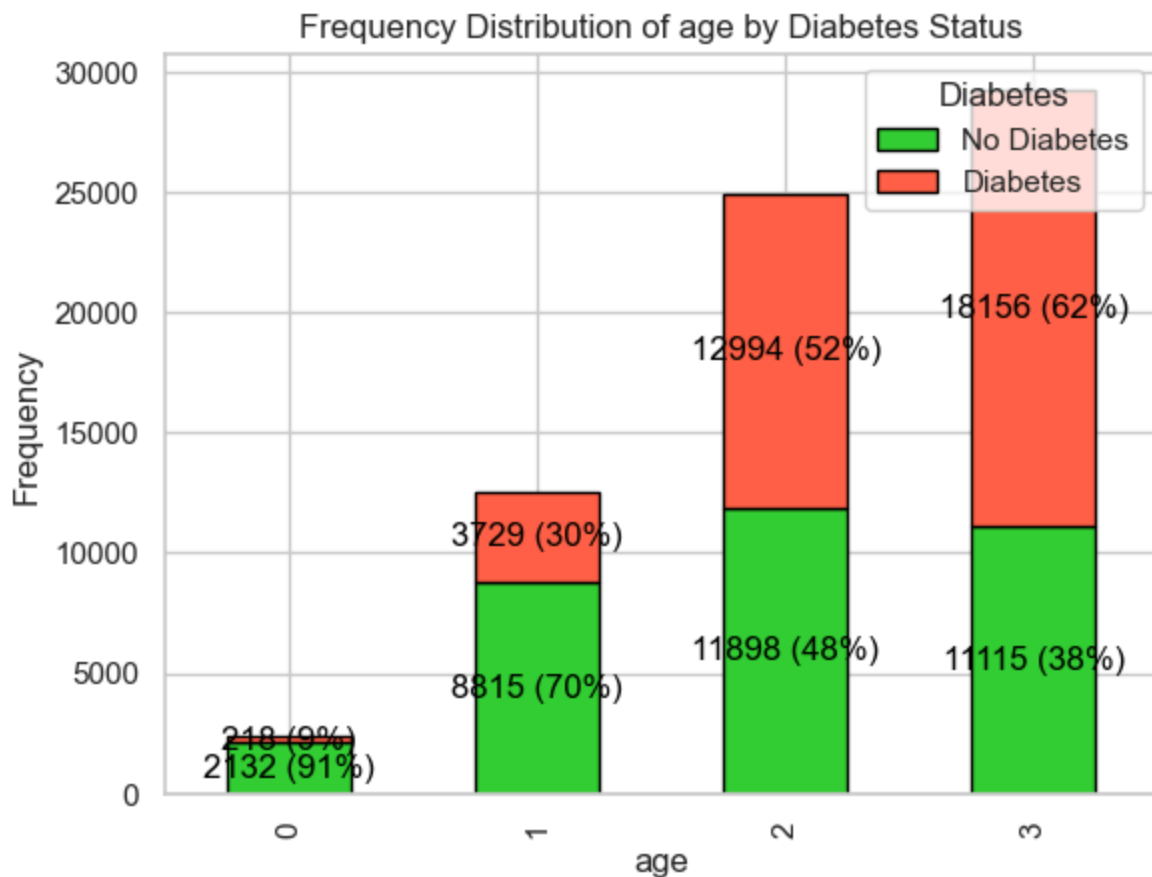
    # Add Labels and title
```

```
plt.xlabel(feature)
plt.ylabel('Frequency')
plt.title(f'Frequency Distribution of {feature} by Diabetes Status')

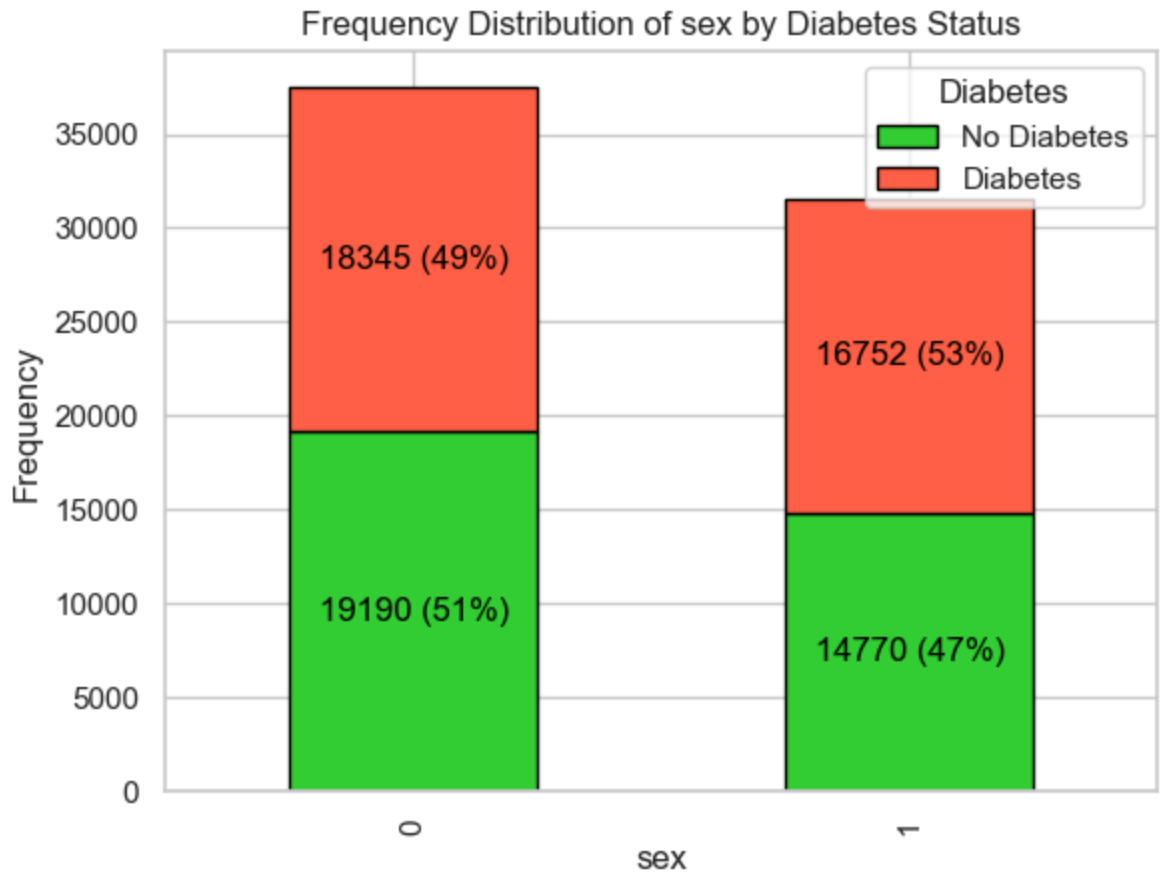
# Add Legend
plt.legend(title='Diabetes', labels=['No Diabetes', 'Diabetes'], loc='upper right')

# Show the plot
plt.show()
```

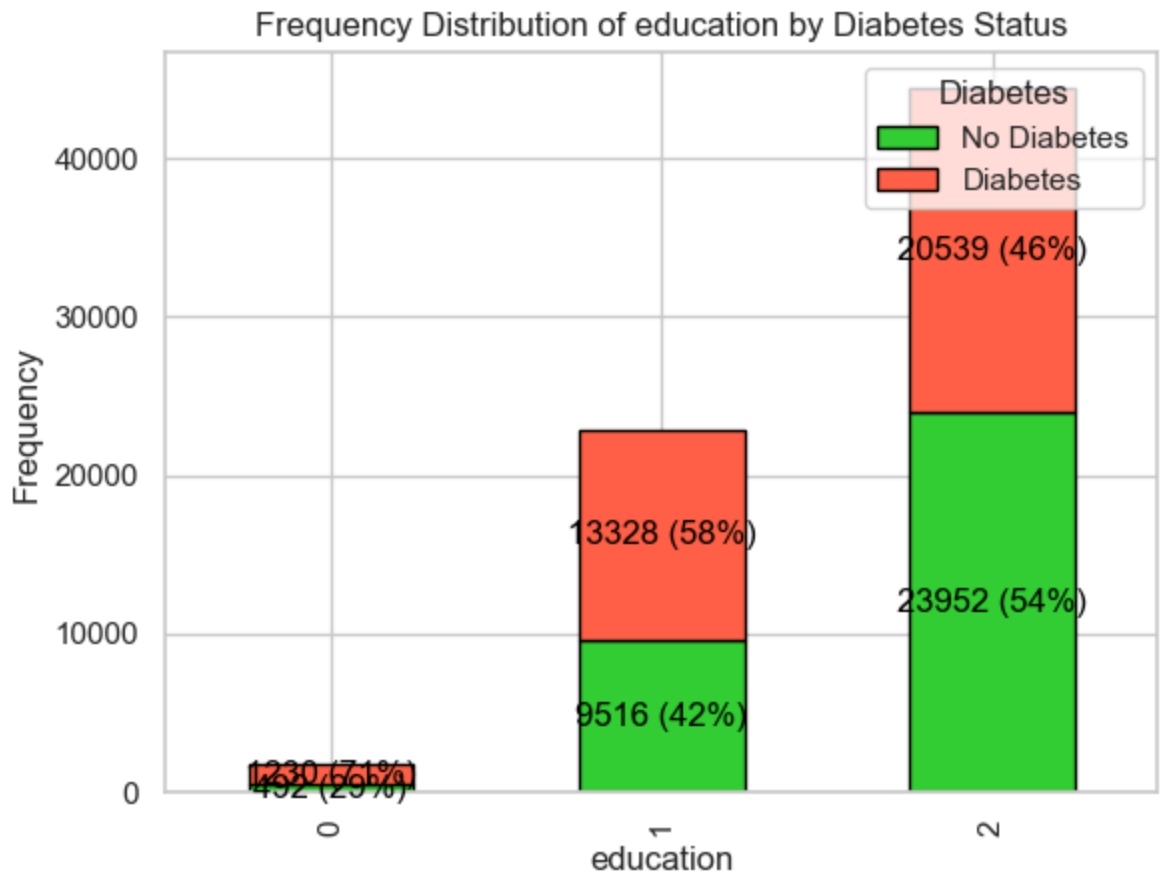
<Figure size 800x400 with 0 Axes>



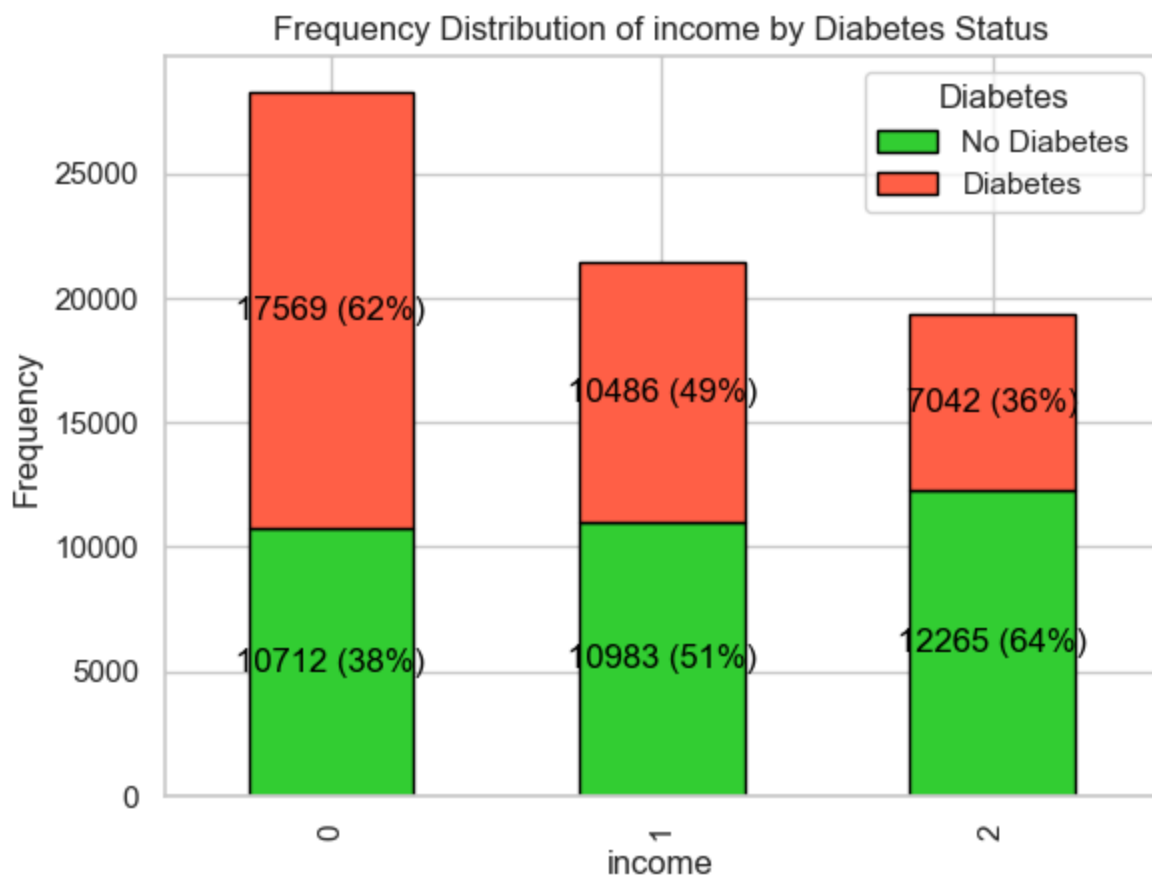
<Figure size 800x400 with 0 Axes>



<Figure size 800x400 with 0 Axes>



<Figure size 800x400 with 0 Axes>



=====

Save the 'df.diabetes' dataset as a separate clean CSV file

- The location of the saved 'diabetes_clean.csv' file will be in the current working directory of my Python environment at the time i run this script

```
In [ ]: import pandas as pd
        from sklearn.datasets import load_diabetes

        # Load diabetes dataset
        diabetes = load_diabetes(as_frame=True)
        df_diabetes = diabetes.frame

        # Save the DataFrame to a CSV file without the index
        df_diabetes.to_csv('diabetes_clean.csv', index=False)
```